Day 6. Soft Actor-Critic (SAC)

NPEX Reinforcement Learning

2020.09.02

Jaeuk Shin



SAC - Review



SAC - Review

How to incentivize exploration?

idea: augment reward as follows:

$$\sum_{t=0}^{T-1} \left(r(s_t, a_t) + \alpha \mathcal{H}(\pi(\cdot|s_t)) \right)$$

where $\mathcal{H}(\pi(\cdot|s_t))$: **entropy** of action distribution $\pi(\cdot|s_t)$

How to solve new MDP with this novel reward criterion?

→ soft Q-learning, soft Bellman equation, etc.





Gaussian actor network $a \sim \mathcal{N}(\mu_{\phi}(s), \sigma_{\phi}(s))$

two critic networks $Q_1(s, a; \theta_1), Q_2(s, a; \theta_2)$

target & loss construction

critic target:

$$y_j = r_j + \gamma \mathbb{E}_{a \sim \pi(\cdot|s_j)} \left(Q(s'_j, a) - \alpha \log \pi(a|s_j) \right)$$

$$\to y_j = r_j + \gamma \mathbb{E}_{a \sim \pi_{\phi^-}(\cdot|s_j)} \left(\min_{i=1,2} Q_i(s'_j, a; \theta_i^-) - \alpha \log \pi_{\phi^-}(a|s_j) \right)$$



actor loss:

$$\alpha \log \pi_{\phi}(f_{\phi}(\epsilon_j, s_j)|s_j) - \min_{i=1,2} Q_i(s_j, f_{\phi}(\epsilon_j, s_j))$$

Remark. π_{ϕ} : probability density, f_{ϕ} : actor network



```
def init (self, dimS, dimA, hidden1, hidden2):
    super(DoubleCritic, self).__init__()
    self.fc1 = nn.Linear(dimS + dimA, hidden1)
    self.fc2 = nn.Linear(hidden1, hidden2)
    self.fc3 = nn.Linear(hidden2, 1)
    self.fc4 = nn.Linear(dimS + dimA, hidden1)
    self.fc5 = nn.Linear(hidden1, hidden2)
    self.fc6 = nn.Linear(hidden2, 1)
def forward(self, state, action):
    x = torch.cat([state, action], dim=1)
   x1 = F.relu(self.fc1(x))
   x1 = F.relu(self.fc2(x1))
   x1 = self.fc3(x1)
   x2 = F.relu(self.fc4(x))
   x2 = F.relu(self.fc5(x2))
   x2 = self.fc6(x2)
    return x1, x2
```

this is how we define twin critics!

```
def Q1(self, state, action):
    x = torch.cat([state, action], dim=1)
    x = F.relu(self.fc1(x))
    x = F.relu(self.fc2(x))
    x = self.fc3(x)
```



```
actor definition - 1<sup>st</sup> step
def forward(self, state, eval=False, with_log_prob=False):
    x = F.relu(self.fc1(state))
    x = F.relu(self.fc2(x))
   mu = self.fc3(x)
    log sigma = self.fc4(x)
    # clip value of log sigma, as was done in Haarnoja's implementation of SAC:
    # https://github.com/haarnoja/sac.git
    log sigma = torch.clamp(log sigma, -20.0, 2.0)
    sigma = torch.exp(log sigma)
    distribution = Independent(Normal(mu, sigma), 1)
```

what we are doing here: compute **distribution params** $\mu_{\phi}(s)$ and $\sigma_{\phi}(s)$



return a, log prob

actor definition - 2nd step

```
if not eval:
   # use rsample() instead of sample(), as sample() does not allow back-propagation through params
   u = distribution.rsample()
                                                                  needed for reparametrization trick
   if with log prob:
       log prob = distribution.log prob(u)
       \log_{prob} = 2.0 * torch.sum((np.log(2.0) + 0.5 * np.log(self.ctrl_range) - u - F.softplus(-2.0 * u)), dim=1)
                                                            tricky part
   else:
                                             u \sim \mathcal{N}(\mu_{\phi}(s), \sigma_{\phi}(s)) \longrightarrow a = \tanh(u) \sim ?
       log prob = None
else:
   u = mu
   log_prob = None
# apply tanh so that the resulting action lies in (-1, 1)^D
a = self.ctrl range * torch.tanh(u)
```

 $softplus(x) = log(1 + e^x)$



self.Q_optimizer.zero_grad()

self.Q_optimizer.step()

Q loss.backward()

```
with torch.no_grad():
    next_actions, log_probs = self.pi(next_obs_batch, with_log_prob=True)
    target q1, target q2 = self.target Q(next obs batch, next actions)
    target q = torch.min(target q1, target q2)
    target = rew batch + self.gamma * masks * (target q - self.alpha * log probs)
out1, out2 = self.Q(obs_batch, act batch)
                                                                                \leftarrow training critics
O loss1 = torch.mean((out1 - target)**2)
Q_loss2 = torch.mean((out2 - target)**2)
                                  trick! (why?)
Q_{loss} = Q_{loss1} + Q_{loss2}
```



```
actions, log_probs = self.pi(obs_batch, with_log_prob=True)
freeze(self.Q)
q1, q2 = self.Q(obs_batch, actions)
q = torch.min(q1, q2)
pi_loss = torch.mean(self.alpha * log_probs - q)
self.pi_optimizer.zero_grad()
pi_loss.backward()
self.pi_optimizer.step()
unfreeze(self.Q)
```

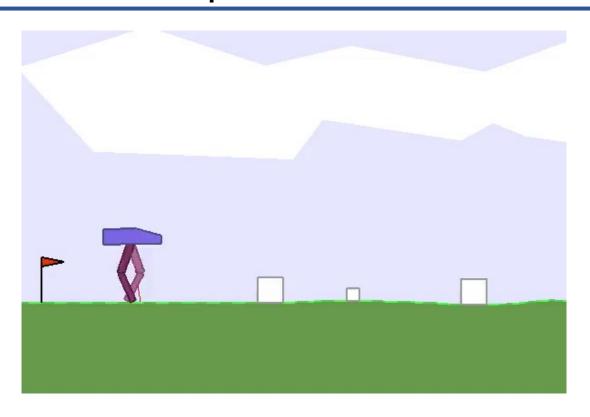
 \leftarrow training actor



SAC - Experiment



SAC - Experiment



Is it successful on BipedalWalker-v3? (state dim: 24 / action dim: 4)





policy gradient (with importance sampling)

$$\nabla_{\phi} J(\phi) = \mathbb{E}_{s \sim \rho_{\phi \text{old}}, \ a \sim \pi_{\phi_{\text{old}}}} \left(A^{\pi_{\phi}}(s, a) \frac{\nabla_{\phi} \pi_{\phi}(a|s)}{\pi_{\phi_{\text{old}}}(a|s)} \right)$$

To estimate $A^{\pi_{\phi}}$, we use a separate value function approximator $V(s;\theta)$, and apply **generalized advantage estimation**(**GAE**)!

Furthermore, we use a stochastic policy $\pi_{\phi}(a|s)$ (Gaussian in our case).

Plus, in TRPO, we have a lot of extra stuff to implement (Hessian-vector product, line search, etc.)



GAE?

Given a trajectory $(s_0, a_0, r_0, s_1, a_1, r_1, \dots s_T)$ generated by executing the current policy, we first compute

$$\delta_t = r_t + \gamma V(s_{t+1}; \theta) - V(s_t; \theta)$$

Then, GAE is computed as follows:

$$\hat{A}(s_t, a_t) = \sum_{\tau=t}^{T-1} \gamma^{\tau-t} \delta_{\tau}$$



training $V(s;\theta)$?

given a trajectory $(s_0, a_0, r_0, s_1, a_1, r_1, \dots s_T)$ generated from π_{ϕ} , we can compute Monte-Carlo esimtates of $V(s_0), V(s_1), \dots V(s_T)$ as follows:

$$V(s_t) = \sum_{\tau=t}^{T-1} \gamma^{\tau-t} r_{\tau} + \gamma^T V(s_T)$$

This is the **target** for value function update!



TRPO - Implementation



TRPO - Implementation

What info do we need when we implement **on-policy algorithms**?

- 1. (s_j, a_j, s'_j, r_j)
- 2. generalized advantage estimation(GAE)
- 3. MC estimates of value $V^{\pi_{\phi}}(s)$
- 4. probability $\pi_{\phi_{\text{old}}}(a_j|s_j)$



TRPO - Implementation

```
self._obs_mem = np.zeros(shape=(lim, dimS))
self._act_mem = np.zeros(shape=(lim, dimA))
self._rew_mem = np.zeros(shape=(lim,))
self._val_mem = np.zeros(shape=(lim,))
self._log_prob_mem = np.zeros(shape=(lim,))
```

collected during agent-env interaction

```
# memory of cumulative rewards which are MC-estimates of the current value function
self._target_v_mem = np.zeros(shape=(lim,))
# memory of GAE($\lambda$)-estimate of the current advantage function
self._adv_mem = np.zeros(shape=(lim,))
```

computed when sampling a single episode is done



Thank you

